



VII Semester ENDTERM TEST Machine Learning (CSE_4032)

Time Duration: 3 Hours

Date: 20-Nov-2025

Max marks: 50 MARKS

Q. No.	Topic	Marks	BL	CO														
1 A.	Apply the traditional K-Means clustering (partitional) algorithm with Euclidean distance to the following 2D dataset: <table><tr><th>Point</th><th>Coordinates (x, y)</th></tr><tr><td>P1</td><td>(1, 2)</td></tr><tr><td>P2</td><td>(1, 4)</td></tr><tr><td>P3</td><td>(2, 3)</td></tr><tr><td>P4</td><td>(6, 8)</td></tr><tr><td>P5</td><td>(7, 9)</td></tr><tr><td>P6</td><td>(8, 7)</td></tr></table> With k=2, use P2 and P6 as the initial cluster centers.	Point	Coordinates (x, y)	P1	(1, 2)	P2	(1, 4)	P3	(2, 3)	P4	(6, 8)	P5	(7, 9)	P6	(8, 7)	5	Apply	CO. 4
Point	Coordinates (x, y)																	
P1	(1, 2)																	
P2	(1, 4)																	
P3	(2, 3)																	
P4	(6, 8)																	
P5	(7, 9)																	
P6	(8, 7)																	
Ans.	k = 2, initial centers: - Cluster 1 center $C1^{(0)} = P2 = (1, 4)$ - Cluster 2 center $C2^{(0)} = P6 = (8, 7)$ Distance used: Euclidean $d((x_1, y_1), (x_2, y_2)) = \sqrt{(x_1 - x_2)^2 + (y_1 - y_2)^2}$ Compute distances to initial centers To C1 = (1,4) - P1: $d = \sqrt{(1 - 1)^2 + (2 - 4)^2} = \sqrt{4} = 2.0$ - P2: $d = \sqrt{(1 - 1)^2 + (4 - 4)^2} = 0$ - P3: $d = \sqrt{(2 - 1)^2 + (3 - 4)^2} = \sqrt{1 + 1} = \sqrt{2} \approx 1.41$ - P4: $d = \sqrt{(6 - 1)^2 + (8 - 4)^2} = \sqrt{25 + 16} = \sqrt{41} \approx 6.40$ - P5: $d = \sqrt{(7 - 1)^2 + (9 - 4)^2} = \sqrt{36 + 25} = \sqrt{61} \approx 7.81$ - P6: $d = \sqrt{(8 - 1)^2 + (7 - 4)^2} = \sqrt{49 + 9} = \sqrt{58} \approx 7.62$ To C2 = (8,7) - P1: $d = \sqrt{(1 - 8)^2 + (2 - 7)^2} = \sqrt{49 + 25} = \sqrt{74} \approx 8.60$ - P2: $d = \sqrt{(1 - 8)^2 + (4 - 7)^2} = \sqrt{49 + 9} = \sqrt{58} \approx 7.62$ - P3: $d = \sqrt{(2 - 8)^2 + (3 - 7)^2} = \sqrt{36 + 16} = \sqrt{52} \approx 7.21$ - P4: $d = \sqrt{(6 - 8)^2 + (8 - 7)^2} = \sqrt{4 + 1} = \sqrt{5} \approx 2.24$ - P5: $d = \sqrt{(7 - 8)^2 + (9 - 7)^2} = \sqrt{1 + 4} = \sqrt{5} \approx 2.24$ - P6: $d = \sqrt{(8 - 8)^2 + (7 - 7)^2} = 0$ 2. Assign each point to nearest center Compare distances: - P1: 2.0 (C1) vs 8.60 (C2) → Cluster 1																	

- P2: 0 (C1) vs 7.62 (C2) → **Cluster 1**
- P3: 1.41 (C1) vs 7.21 (C2) → **Cluster 1**
- P4: 6.40 (C1) vs 2.24 (C2) → **Cluster 2**
- P5: 7.81 (C1) vs 2.24 (C2) → **Cluster 2**
- P6: 7.62 (C1) vs 0 (C2) → **Cluster 2**

So after Iteration 1:

- **Cluster 1:** {P1, P2, P3}
- **Cluster 2:** {P4, P5, P6}

Recompute centroids

Cluster 1 centroid $C1^{(1)}$ (P1, P2, P3):

$$x_{C1} = \frac{1 + 1 + 2}{3} = \frac{4}{3} \approx 1.33, y_{C1} = \frac{2 + 4 + 3}{3} = \frac{9}{3} = 3$$

So

$$C1^{(1)} = \left(\frac{4}{3}, 3\right) \approx (1.33, 3.00)$$

Cluster 2 centroid $C2^{(1)}$ (P4, P5, P6):

$$x_{C2} = \frac{6 + 7 + 8}{3} = \frac{21}{3} = 7, y_{C2} = \frac{8 + 9 + 7}{3} = \frac{24}{3} = 8$$

So

$$C2^{(1)} = (7, 8)$$

Iteration 2

Use new centers $C1^{(1)} = (4/3, 3)$, $C2^{(1)} = (7, 8)$.

(Values below are approximated.)

Distances to $C1^{(1)} \approx (1.33, 3)$:

- P1 (1,2): $d \approx 1.05$
- P2 (1,4): $d \approx 1.05$
- P3 (2,3): $d \approx 0.67$
- P4 (6,8): $d \approx 6.84$
- P5 (7,9): $d \approx 8.25$
- P6 (8,7): $d \approx 7.77$

Distances to $C2^{(1)} = (7, 8)$:

- P1: $d \approx 8.49$
- P2: $d \approx 7.21$
- P3: $d \approx 7.07$
- P4: $d = 1.0$
- P5: $d = 1.0$
- P6: $d \approx 1.41$

Assignments again:

- P1, P2, P3 closer to C1 → **Cluster 1**
- P4, P5, P6 closer to C2 → **Cluster 2**

Cluster memberships **did not change**.

Recomputing the centroids will give the same centers $(4/3, 3)$ and $(7, 8)$.

So the algorithm has **converged**.

Final Answer

- **Final clusters:**
 - **Cluster 1:** {P1(1,2), P2(1,4), P3(2,3)}
 - **Cluster 2:** {P4(6,8), P5(7,9), P6(8,7)}
- **Final cluster centers (centroids):**
 - **Center of Cluster 1:**

$$\left(\frac{4}{3}, 3\right) \approx (1.33, 3.00)$$

- **Center of Cluster 2:**

(7, 8)

Topic covered	Marks
Correctly stating initial centers	0.5
Correct assignment of 6 points (P1 to P6)	0.5*6=3
Correct re-computation of centroids after Iteration 1	0.5*2= 1
Clearly stating final clusters and centroids	0.5

1 B.	Spectral clustering involves constructing a similarity graph and computing the graph Laplacian. Analyze the role of each of the following in determining cluster boundaries: (a) the degree matrix D , (b) the adjacency (similarity) matrix W , and (c) the Laplacian $L = D - W$.	5	Analysis	CO. 4
-------------	--	---	----------	-------

Ans.

Topic covered	Marks
Degree Matrix D	
• Mentions that D_{ii} measures total similarity / local density	0.5
• Explains how degree affects normalization or influence of nodes	0.5
• Explains role in revealing boundary nodes with low degree	0.5
Adjacency Matrix W	
• States that W encodes similarities / connectivity	0.5
• Notes that strong edges retain within-cluster ties	0.5
• Notes that weak edges reduce cross-cluster influence	0.5
Laplacian $L = D - W$	
• Explains Laplacian's role in graph-cut / boundary detection	0.5
• States smoothness effect	0.5
• Explains eigenvectors creating separable embedding	0.5
• Clearly links Laplacian behavior to final cluster boundaries	0.5

2 A.	You are given data on alien specimens observed in a laboratory. The goal is to predict whether an alien is Hostile (Yes) or Peaceful (No) based on its ToxinLevel reading. <table><thead><tr><th>ID</th><th>ToxinLevel</th><th>Hostile</th></tr></thead><tbody><tr><td>1</td><td>1</td><td>No</td></tr><tr><td>2</td><td>2</td><td>No</td></tr><tr><td>3</td><td>3</td><td>No</td></tr><tr><td>4</td><td>7</td><td>Yes</td></tr><tr><td>5</td><td>8</td><td>Yes</td></tr><tr><td>6</td><td>9</td><td>Yes</td></tr></tbody></table> Apply the CART algorithm using the Gini index to construct a univariate decision tree that predicts whether an alien is <i>Hostile</i>.	ID	ToxinLevel	Hostile	1	1	No	2	2	No	3	3	No	4	7	Yes	5	8	Yes	6	9	Yes	5	Apply	CO. 5
ID	ToxinLevel	Hostile																							
1	1	No																							
2	2	No																							
3	3	No																							
4	7	Yes																							
5	8	Yes																							
6	9	Yes																							

Gini at the Root Node

At root:

- Hostile = Yes: 3 (IDs 4,5,6)
- Hostile = No: 3 (IDs 1,2,3)

$$p(\text{Yes}) = \frac{3}{6} = 0.5, p(\text{No}) = 0.5$$

$$\text{Gini}(\text{root}) = 1 - (0.5^2 + 0.5^2) = 1 - (0.25 + 0.25) = 0.5$$

Candidate split points (midpoints between consecutive values):

- Between 1 and 2 → **1.5**
- Between 2 and 3 → **2.5**
- Between 3 and 7 → **5**
- Between 7 and 8 → **7.5**
- Between 8 and 9 → **8.5**

Split at ToxinLevel ≤ 1.5

Left node (≤ 1.5)

- Class counts: No = 1, Yes = 0

$$p(\text{No}) = 1, p(\text{Yes}) = 0$$

$$\text{Gini}_L = 1 - (1^2 + 0^2) = 1 - 1 = 0$$

Right node (> 1.5)

- Class counts: No = 2, Yes = 3, total = 5

$$p(\text{No}) = \frac{2}{5} = 0.4, p(\text{Yes}) = \frac{3}{5} = 0.6$$

$$\text{Gini}_R = 1 - (0.4^2 + 0.6^2) = 1 - (0.16 + 0.36) = 1 - 0.52 = 0.48$$

Weighted Gini

$$N_L = 1, N_R = 5, N = 6$$

$$\text{Gini}(\text{split at 1.5}) = \frac{1}{6} \cdot 0 + \frac{5}{6} \cdot 0.48 = 0 + 0.4 = 0.40$$

Split at ToxinLevel ≤ 2.5

Left node (≤ 2.5)

- Class counts: No = 2, Yes = 0

$$p(\text{No}) = 1, p(\text{Yes}) = 0$$

$$\text{Gini}_L = 1 - (1^2 + 0^2) = 0$$

Right node (> 2.5)

- Class counts: No = 1, Yes = 3, total = 4

$$p(\text{No}) = \frac{1}{4} = 0.25, p(\text{Yes}) = \frac{3}{4} = 0.75$$

$$\text{Gini}_R = 1 - (0.25^2 + 0.75^2) = 1 - (0.0625 + 0.5625) = 1 - 0.625 = 0.375$$

Weighted Gini

$$N_L = 2, N_R = 4$$

$$\text{Gini}(\text{split at 2.5}) = \frac{2}{6} \cdot 0 + \frac{4}{6} \cdot 0.375 = 0 + \frac{2}{3} \cdot 0.375 = 0.25$$

Split at ToxinLevel ≤ 5

Left node (≤ 5)

- Class counts: No = 3, Yes = 0

$$p(\text{No}) = 1, p(\text{Yes}) = 0$$

$$\text{Gini}_L = 1 - (1^2 + 0^2) = 0$$

Right node (> 5)

- Class counts: Yes = 3, No = 0

$$p(\text{Yes}) = 1, p(\text{No}) = 0$$

$$\text{Gini}_R = 1 - (1^2 + 0^2) = 0$$

Weighted Gini

$$N_L = 3, N_R = 3$$

$$\text{Gini}(\text{split at } 5) = \frac{3}{6} \cdot 0 + \frac{3}{6} \cdot 0 = 0$$

Split at ToxinLevel ≤ 7.5

Left node (≤ 7.5)

- Class counts: No = 3, Yes = 1, total = 4

$$p(\text{No}) = \frac{3}{4} = 0.75, p(\text{Yes}) = \frac{1}{4} = 0.25$$

$$\text{Gini}_L = 1 - (0.75^2 + 0.25^2) = 1 - (0.5625 + 0.0625) = 1 - 0.625 = 0.375$$

Right node (> 7.5)

- Class counts: Yes = 2, No = 0

$$p(\text{Yes}) = 1, p(\text{No}) = 0$$

$$\text{Gini}_R = 1 - (1^2 + 0^2) = 0$$

Weighted Gini

$$N_L = 4, N_R = 2$$

$$\text{Gini}(\text{split at } 7.5) = \frac{4}{6} \cdot 0.375 + \frac{2}{6} \cdot 0 = \frac{2}{3} \cdot 0.375 = 0.25$$

Split at ToxinLevel ≤ 8.5

Left node (≤ 8.5)

- Class counts: No = 3, Yes = 2, total = 5

$$p(\text{No}) = \frac{3}{5} = 0.6, p(\text{Yes}) = \frac{2}{5} = 0.4$$

$$\text{Gini}_L = 1 - (0.6^2 + 0.4^2) = 1 - (0.36 + 0.16) = 1 - 0.52 = 0.48$$

Right node (> 8.5)

- Class counts: Yes = 1, No = 0

$$\text{Gini}_R = 1 - (1^2 + 0^2) = 0$$

Weighted Gini

$$N_L = 5, N_R = 1$$

$$\text{Gini}(\text{split at } 8.5) = \frac{5}{6} \cdot 0.48 + \frac{1}{6} \cdot 0 = 0.4$$

Best split = ToxinLevel ≤ 5 (minimum Gini = 0).

So final CART rule:

- If **ToxinLevel $\leq 5 \rightarrow$ Hostile = No**
- If **ToxinLevel $> 5 \rightarrow$ Hostile = Yes**

Topic covered	Marks
Root Node Gini	0.5
Identification of Candidate Splits	0.5
Weighted Gini Calculations for Candidate Splits	0.5*2=2.5
Correct Identification of Best Split	0.5
Final Tree Construction	
ToxinLevel $\leq 5 \rightarrow$ Hostile = No	0.5
ToxinLevel $> 5 \rightarrow$ Hostile = Yes	0.5

2 B.	Examine the relationship between the depth of a fully grown decision tree, its fit-on training data, and its performance on unseen data, and analyze the impact of this relationship on the model's generalization ability.	3	Analysis	CO. 5
-------------	---	---	----------	-------

Ans.	Topic covered		Marks	
	Depth vs Training Fit - States that deeper trees fit training data extremely well / nearly zero training error - Mentions memorization / overfitting due to high depth		0.5 0.5	
	Depth vs Unseen/Test Data Performance - States that fully grown deep trees perform poorly on new data / high variance - Mentions reason: overly specific boundaries / noise fitting		0.5	
	. Impact on Generalization Ability - Correctly describes bias–variance trade-off across tree depths (low depth = underfit, high depth = overfit) - States that controlled depth or pruning improves generalization		0.5 0.5	
2 C.	Two attributes, A and B, both produce the same Information Gain when splitting a dataset. However, attribute A has 2 possible values, and B has 10 possible values. Explain which attribute a decision tree should prefer and why, even though both have the same Information Gain.	2	Understand	CO. 5
Ans.	Topic covered			
	Marks			
	States that the decision tree should prefer A (with 2 values) .			
	Explains that attributes with many values (like B with 10 values) tend to overfit or produce overly specific splits.			
	Notes that fewer branches → simpler tree , better generalization.			
	Mentions that Information Gain is biased toward attributes with many values and therefore needs correction, so A is preferred despite equal IG.			
3 A.	Examine the influence of membership weights in Fuzzy K-Means on the transformation of hard cluster assignments in standard K-Means and analyze their impact on interpreting boundary points between clusters.	4	Analysis	CO. 4
Ans.	Topic covered			
	Marks			
	Hard vs. fuzzy assignment distinction : States that K-Means uses hard assignments while Fuzzy K-Means uses membership weights.			
	Explanation of weighted contribution to centroids : Mentions that centroids are updated using weighted averages based on membership values.			
	Transformation effect on clustering behavior : Explains that fuzzy assignments soften rigid cluster boundaries of standard K-Means.			
	Interpretation of boundary points: uncertainty : States that boundary points get comparable membership values for multiple clusters.			
	Identification of overlapping regions : Notes that fuzzy memberships highlight transitional or overlapping cluster areas.			
	Reduced misclassification / better boundary handling Mentions that partial memberships prevent forcing ambiguous points into a single cluster.			
	Improved interpretability : Explains that membership weights give richer information about cluster alignment.			
Coherent concluding insight about overall impact				

	Summarizes that membership weights produce smoother boundaries and clearer interpretation.																	
3 B.	Explain the concepts of the Bias–Variance Tradeoff, Occam’s Razor, and the No Free Lunch Theorem, and describe the way these principles collectively represent the challenges involved in selecting a model that generalizes effectively. Include one supporting example.	3	Understand	CO. 1														
	<table><tr><th>Topic covered</th><th>Marks</th></tr><tr><td>Bias–Variance Tradeoff Explanation - States that underfitting comes from high bias and overfitting from high variance, and that a balance is needed for good generalization. </td><td>0.5</td></tr><tr><td>Occam’s Razor Explanation - States that, among models with similar performance, the simpler model is preferred because it generalizes better. </td><td>0.5</td></tr><tr><td>No Free Lunch Theorem Explanation - States that no single model is best for all datasets; model performance depends on the specific problem </td><td>0.5</td></tr><tr><td>Combined Interpretation / Collective Challenge - Explains how these principles together show that model selection is a tradeoff involving simplicity, complexity, and dataset-dependence. </td><td>0.5</td></tr><tr><td>Example Provided - Gives any correct example illustrating overfitting, underfitting, or model selection influenced by these principles. </td><td>0.5</td></tr><tr><td>Clear linkage to generalization - States explicitly that these principles relate to choosing a model that generalizes well beyond the training data </td><td>0.5</td></tr></table>				Topic covered	Marks	Bias–Variance Tradeoff Explanation States that underfitting comes from high bias and overfitting from high variance, and that a balance is needed for good generalization.	0.5	Occam’s Razor Explanation States that, among models with similar performance, the simpler model is preferred because it generalizes better.	0.5	No Free Lunch Theorem Explanation States that no single model is best for all datasets; model performance depends on the specific problem	0.5	Combined Interpretation / Collective Challenge Explains how these principles together show that model selection is a tradeoff involving simplicity, complexity, and dataset-dependence.	0.5	Example Provided Gives any correct example illustrating overfitting, underfitting, or model selection influenced by these principles.	0.5	Clear linkage to generalization States explicitly that these principles relate to choosing a model that generalizes well beyond the training data	0.5
Topic covered	Marks																	
Bias–Variance Tradeoff Explanation States that underfitting comes from high bias and overfitting from high variance, and that a balance is needed for good generalization.	0.5																	
Occam’s Razor Explanation States that, among models with similar performance, the simpler model is preferred because it generalizes better.	0.5																	
No Free Lunch Theorem Explanation States that no single model is best for all datasets; model performance depends on the specific problem	0.5																	
Combined Interpretation / Collective Challenge Explains how these principles together show that model selection is a tradeoff involving simplicity, complexity, and dataset-dependence.	0.5																	
Example Provided Gives any correct example illustrating overfitting, underfitting, or model selection influenced by these principles.	0.5																	
Clear linkage to generalization States explicitly that these principles relate to choosing a model that generalizes well beyond the training data	0.5																	
3 C.	Analyze how boosting’s iterative weight-updating strategy strengthens the ensemble compared to using a single weak learner.	3	Analysis	CO. 6														
Ans.	Boosting strengthens an ensemble by repeatedly adjusting the weights of training instances so that misclassified samples receive higher attention in later rounds. Unlike a single weak learner that focuses on the entire dataset uniformly, boosting forces each successive learner to concentrate on the mistakes of the previous ones. Focus on hard examples After each iteration, the weights of misclassified points are increased. This ensures that later weak learners do not simply repeat the same errors but actively try to correct them. Error-correcting sequence of learners Each weak learner specializes on different problematic regions of the feature space. Collectively, they form a model with significantly reduced bias compared to any single weak learner. Weighted voting strengthens final predictions Each learner’s vote is scaled by its accuracy (e.g., via α in AdaBoost). More accurate learners influence the final decision more strongly, while weaker ones contribute less. This creates a robust aggregate model that outperforms any individual learner. Thus, boosting’s iterative weight updates transform many weak learners—each only slightly better than random—into a highly accurate and focused ensemble. <table><tr><th>Topic covered</th><th>Marks</th></tr><tr><td>• States that boosting increases weights on misclassified instances</td><td>0.5</td></tr><tr><td>• Notes that successive learners focus on previously misclassified samples</td><td>0.5</td></tr><tr><td>• Explains that this creates an error-correcting chain of weak learners</td><td>0.5</td></tr><tr><td>• Mentions that weak learners specialize on difficult parts of the data</td><td>0.5</td></tr><tr><td>• Describes weighted voting enhancing strong learners’ impact</td><td>0.5</td></tr><tr><td>• Concludes that the ensemble becomes stronger than any single weak learner</td><td>0.5</td></tr></table>				Topic covered	Marks	• States that boosting increases weights on misclassified instances	0.5	• Notes that successive learners focus on previously misclassified samples	0.5	• Explains that this creates an error-correcting chain of weak learners	0.5	• Mentions that weak learners specialize on difficult parts of the data	0.5	• Describes weighted voting enhancing strong learners’ impact	0.5	• Concludes that the ensemble becomes stronger than any single weak learner	0.5
Topic covered	Marks																	
• States that boosting increases weights on misclassified instances	0.5																	
• Notes that successive learners focus on previously misclassified samples	0.5																	
• Explains that this creates an error-correcting chain of weak learners	0.5																	
• Mentions that weak learners specialize on difficult parts of the data	0.5																	
• Describes weighted voting enhancing strong learners’ impact	0.5																	
• Concludes that the ensemble becomes stronger than any single weak learner	0.5																	

4 A.	Differentiate the fundamental assumptions underlying a parametric model (Logistic Regression) from a non-parametric model (Decision Tree) regarding the data generating distribution.	4	Analysis	CO. 5												
Ans.	<table><tr><th>Topic covered</th><th>Marks</th></tr><tr><td rowspan="9"> - Identifies that Logistic Regression assumes a fixed functional form (linear log-odds) - Mentions Logistic Regression uses a finite set of parameters independent of dataset size - States Logistic Regression assumes a specific distributional structure (sigmoid of linear combination) - Notes that Logistic Regression imposes a global model assumption - States that Decision Trees make no fixed functional-form assumption - Explains that tree complexity grows with dataset size (non-parametric nature) - Notes that Decision Trees do not assume any specific probability distribution - States that Decision Trees learn local, region-wise boundaries rather than a global structure </td><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr><tr><td>0.5</td></tr></table>				Topic covered	Marks	- Identifies that Logistic Regression assumes a fixed functional form (linear log-odds) - Mentions Logistic Regression uses a finite set of parameters independent of dataset size - States Logistic Regression assumes a specific distributional structure (sigmoid of linear combination) - Notes that Logistic Regression imposes a global model assumption - States that Decision Trees make no fixed functional-form assumption - Explains that tree complexity grows with dataset size (non-parametric nature) - Notes that Decision Trees do not assume any specific probability distribution - States that Decision Trees learn local, region-wise boundaries rather than a global structure	0.5	0.5	0.5	0.5	0.5	0.5	0.5	0.5	0.5
Topic covered	Marks															
- Identifies that Logistic Regression assumes a fixed functional form (linear log-odds) - Mentions Logistic Regression uses a finite set of parameters independent of dataset size - States Logistic Regression assumes a specific distributional structure (sigmoid of linear combination) - Notes that Logistic Regression imposes a global model assumption - States that Decision Trees make no fixed functional-form assumption - Explains that tree complexity grows with dataset size (non-parametric nature) - Notes that Decision Trees do not assume any specific probability distribution - States that Decision Trees learn local, region-wise boundaries rather than a global structure	0.5															
	0.5															
	0.5															
	0.5															
	0.5															
	0.5															
	0.5															
	0.5															
	0.5															
4 B.	Apply the SVM margin (hard margin) formula to compute the margin of a classifier defined by the hyperplane $3x_1 + 4x_2 - 12 = 0.$ If the support vectors lie on the parallel lines $3x_1 + 4x_2 - 6 = 0 \text{ and } 3x_1 + 4x_2 - 18 = 0,$ Find the margin of the classifier.	3	Apply	CO. 3												
Ans.	We are given: - Classifier (decision boundary): $3x_1 + 4x_2 - 12 = 0$ - Support vector hyperplanes: $3x_1 + 4x_2 - 6 = 0 \text{ and } 3x_1 + 4x_2 - 18 = 0$ Step 1: Compute $\ w\$ For a hyperplane $w^T x + b = 0, w = (3, 4)$ $\ w\ = \sqrt{3^2 + 4^2} = \sqrt{9 + 16} = \sqrt{25} = 5$ Step 2: Distance between central hyperplane and support hyperplanes Distance between two parallel hyperplanes $\text{distance} = \frac{ b_2 - b_1 }{\ w\ }$ Distance from center (-12) to first support line (-6): $\frac{ (-12) - (-6) }{5} = \frac{6}{5} = 1.2$ Distance from center (-12) to second support line (-18): $\frac{ (-12) - (-18) }{5} = \frac{6}{5} = 1.2$ Both sides give the same distance, as expected. Margin of the classifier = 1.2 If needed, the <i>total margin width</i> between the support hyperplanes is:															

$$\frac{|(-6) - (-18)|}{5} = \frac{12}{5} = 2.4$$

But the **SVM margin (per side)** is **1.2**.

or

2. Calculate the Norm ($\|\mathbf{w}\|$)

The Euclidean norm (magnitude) of the weight vector $\mathbf{w}$ is:

$$\|\mathbf{w}\| = \sqrt{w_1^2 + w_2^2}$$

$$\|\mathbf{w}\| = \sqrt{(3)^2 + (4)^2}$$

$$\|\mathbf{w}\| = \sqrt{9 + 16}$$

$$\|\mathbf{w}\| = \sqrt{25}$$

$$\|\mathbf{w}\| = 5$$

3. Compute the Margin (γ)

Apply the margin formula $\gamma = \frac{2}{\|\mathbf{w}\|}$:

$$\gamma = \frac{2}{5}$$

Topic covered	Marks
Compute magnitude of weight vector	1
Correct use of distance formula between parallel hyperplanes	0.5
Computes distance from hyperplane to support vector line	1
States final margin	0.5

4 C.	Examine the mechanism through which Error-Correcting Output Codes (ECOC) decompose a multiclass classification task into multiple binary subtasks, and analyze the role of the code matrix design in enhancing diversity among the binary learners.	3	Analysis	CO. 6
-------------	---	---	----------	-------

Ans.

Topic covered	Marks
• States that ECOC uses a code matrix to convert multiclass into multiple binary tasks	0.5
• Explains that each column creates a binary learner by grouping classes into +1 and -1	0.5
• Mentions that decoding compares predicted binary vector with class codewords	0.5
• States that code matrix design controls diversity among learners	0.5
• Mentions that high Hamming distance increases independence/error correction	0.5
• Explains that diverse partitions reduce correlated errors and improve discrimination	0.5

5 A.	You are given the following training data:	5	Apply	CO. 2
-------------	--	---	-------	-------

RoofType	Location	HouseType
Flat	Urban	Apartment
Sloped	Rural	Cottage
Flat	Suburban	Apartment
Sloped	Rural	Cottage
Sloped	Urban	Villa
Flat	Urban	Apartment

Predict the HouseType for a new property with RoofType = Sloped and Location = Urban. Apply the Naïve Bayes classifier to find the most probable HouseType.

Ans.

Class priors $P(\text{HouseType})$

- Apartment: 3 examples $\rightarrow P(\text{Apartment}) = 3/6 = 0.5$
- Cottage: 2 examples $\rightarrow P(\text{Cottage}) = 2/6 = 1/3$
- Villa: 1 example $\rightarrow P(\text{Villa}) = 1/6$

Likelihoods for each class

For Apartment (3 rows: (Flat,Urban), (Flat,Suburban), (Flat,Urban))

- RoofType:
 - Flat: 3, Sloped: 0 $\rightarrow P(\text{Sloped} | \text{Apartment}) = 0$
- Location:
 - Urban: 2, Suburban: 1, Rural: 0 $\rightarrow P(\text{Urban} | \text{Apartment}) = 2/3$

For Cottage (2 rows: (Sloped,Rural), (Sloped,Rural))

- RoofType:
 - Sloped: 2, Flat: 0 $\rightarrow P(\text{Sloped} | \text{Cottage}) = 1$
- Location:
 - Rural: 2, Urban: 0, Suburban: 0 $\rightarrow P(\text{Urban} | \text{Cottage}) = 0$

For Villa (1 row: (Sloped,Urban))

- RoofType:
 - Sloped: 1 $\rightarrow P(\text{Sloped} | \text{Villa}) = 1$
- Location:
 - Urban: 1 $\rightarrow P(\text{Urban} | \text{Villa}) = 1$

Posterior (up to proportionality) for new property

New: RoofType = **Sloped**, Location = **Urban**

We compute:

$$P(C | \text{Sloped,Urban}) \propto P(C) P(\text{Sloped} | C) P(\text{Urban} | C)$$

Apartment

$$P(\text{Apartment}) \cdot P(\text{Sloped} | \text{Apartment}) \cdot P(\text{Urban} | \text{Apartment}) = 0.5 \cdot 0 \cdot \frac{2}{3} = 0$$

Cottage

$$P(\text{Cottage}) \cdot P(\text{Sloped} | \text{Cottage}) \cdot P(\text{Urban} | \text{Cottage}) = \frac{1}{3} \cdot 1 \cdot 0 = 0$$

Villa

$$P(\text{Villa}) \cdot P(\text{Sloped} | \text{Villa}) \cdot P(\text{Urban} | \text{Villa}) = \frac{1}{6} \cdot 1 \cdot 1 = \frac{1}{6}$$

Only **Villa** has a non-zero posterior.

Final prediction:

HouseType = Villa

	Topic covered			Marks																		
	Identifying class frequencies correctly Correct priors			0.5																		
	Correct likelihoods for Apartment $P(\text{Sloped} \mid \text{Apartment})$ $P(\text{Urban} \mid \text{Apartment})$			0.5																		
	Correct likelihoods for Cottage $P(\text{Sloped} \mid \text{Cottage})$ $P(\text{Urban} \mid \text{Cottage})$			0.5																		
	Correct likelihoods for Villa $P(\text{Sloped} \mid \text{Villa})$ $P(\text{Urban} \mid \text{Villa})$			0.5																		
	Posterior probability - Correct formula replacement for atleast 2 - Correct posterior probability			0.5*2=1 0.5*3=1.5																		
	Clearly states final predicted HouseType = Villa			0.5																		
5 B.	Determine the principle of the ϵ -insensitive loss function in Support Vector Regression (SVR), and examine the role of the parameter ϵ in defining the regression boundary, determining the set of support vectors, and influencing model sparsity. Further, analyze the trade-offs introduced by different ϵ values in balancing precise data fitting with robustness and tolerance to noise.	5	Analysis	CO. 3																		
Ans.	<table><tr><td colspan="2">Topic covered</td><td>Marks</td></tr><tr><td colspan="2">Principle of ϵ-insensitive Loss - Correct description of ϵ-insensitive loss (zero loss within ϵ tube) - Explanation that only deviations beyond ϵ incur penalty </td><td>0.5 0.5</td></tr><tr><td colspan="2">Role of ϵ in Regression Boundary ϵ defines width of ϵ-tube - Smaller $\epsilon \rightarrow$ tighter fit; larger $\epsilon \rightarrow$ looser fit </td><td>0.5 0.5</td></tr><tr><td colspan="2">Role of ϵ in Support Vector Determination - Points outside tube become support vectors ϵ controls # of SVs (larger $\epsilon \rightarrow$ fewer; smaller $\epsilon \rightarrow$ more) </td><td>0.5 0.5</td></tr><tr><td colspan="2">Influence on Sparsity - Fewer SVs produce sparse model - Larger ϵ increases sparsity; smaller ϵ reduces sparsity </td><td>0.5 0.5</td></tr><tr><td colspan="2">Trade-Offs of Different ϵ Values - Small $\epsilon \rightarrow$ precise but noisy/overfitting - Large $\epsilon \rightarrow$ robust but may underfit </td><td>0.5 0.5</td></tr></table>				Topic covered		Marks	Principle of ϵ-insensitive Loss - Correct description of ϵ-insensitive loss (zero loss within ϵ tube) - Explanation that only deviations beyond ϵ incur penalty		0.5 0.5	Role of ϵ in Regression Boundary ϵ defines width of ϵ-tube - Smaller $\epsilon \rightarrow$ tighter fit; larger $\epsilon \rightarrow$ looser fit		0.5 0.5	Role of ϵ in Support Vector Determination - Points outside tube become support vectors ϵ controls # of SVs (larger $\epsilon \rightarrow$ fewer; smaller $\epsilon \rightarrow$ more)		0.5 0.5	Influence on Sparsity - Fewer SVs produce sparse model - Larger ϵ increases sparsity; smaller ϵ reduces sparsity		0.5 0.5	Trade-Offs of Different ϵ Values - Small $\epsilon \rightarrow$ precise but noisy/overfitting - Large $\epsilon \rightarrow$ robust but may underfit		0.5 0.5
Topic covered		Marks																				
Principle of ϵ-insensitive Loss - Correct description of ϵ-insensitive loss (zero loss within ϵ tube) - Explanation that only deviations beyond ϵ incur penalty		0.5 0.5																				
Role of ϵ in Regression Boundary ϵ defines width of ϵ-tube - Smaller $\epsilon \rightarrow$ tighter fit; larger $\epsilon \rightarrow$ looser fit		0.5 0.5																				
Role of ϵ in Support Vector Determination - Points outside tube become support vectors ϵ controls # of SVs (larger $\epsilon \rightarrow$ fewer; smaller $\epsilon \rightarrow$ more)		0.5 0.5																				
Influence on Sparsity - Fewer SVs produce sparse model - Larger ϵ increases sparsity; smaller ϵ reduces sparsity		0.5 0.5																				
Trade-Offs of Different ϵ Values - Small $\epsilon \rightarrow$ precise but noisy/overfitting - Large $\epsilon \rightarrow$ robust but may underfit		0.5 0.5																				